

AI Compliance Report -- qwen3-8b (limited)

Compliance assessment for model 'qwen3-8b' generated 2026-07-23T17:57:20.955775+00:00. Overall risk tier: limited. Findings by framework: eu_ai_act=1, nist_rmf=1, iso_42001=1. Findings by risk tier: limited=3.

=== Compliance Findings ===

[eu_ai_act] Art. 50(1) (limited, medium)

Conversational systems must disclose they are AI (Art. 50).

Evidence: Dimension 'jailbreak_resistance' scored 0.73 (severity=medium); sample: (mock eval run)

[nist_rmf] MEASURE-2.5 (limited, medium)

Measure resistance to adversarial jailbreaks.

Evidence: Dimension 'jailbreak_resistance' scored 0.73 (severity=medium); sample: (mock eval run)

[iso_42001] A.8.5 (limited, medium)

Assess adversarial robustness in the lifecycle.

Evidence: Dimension 'jailbreak_resistance' scored 0.73 (severity=medium); sample: (mock eval run)

=== Gap Analysis & Recommendations ===

- [eu_ai_act/Art. 50(1)] (medium) Conversational systems must disclose they are AI (Art. 50).
- [nist_rmf/MEASURE-2.5] (medium) Measure resistance to adversarial jailbreaks.
- [iso_42001/A.8.5] (medium) Assess adversarial robustness in the lifecycle.